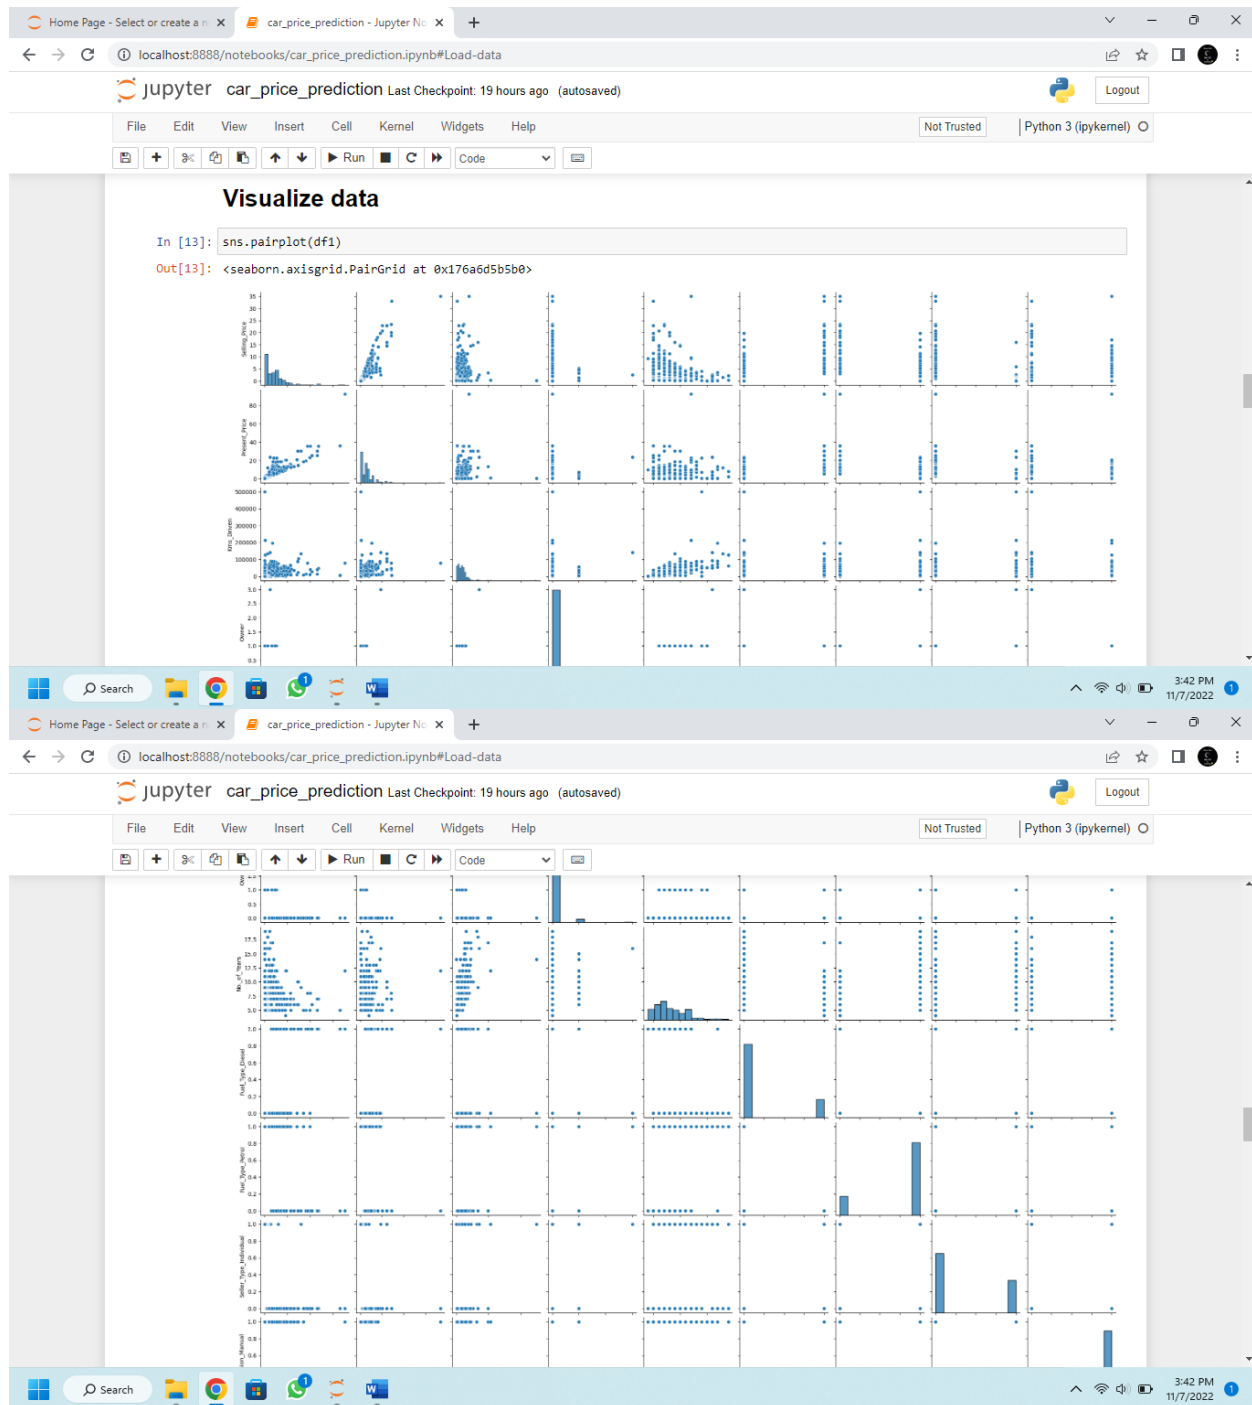
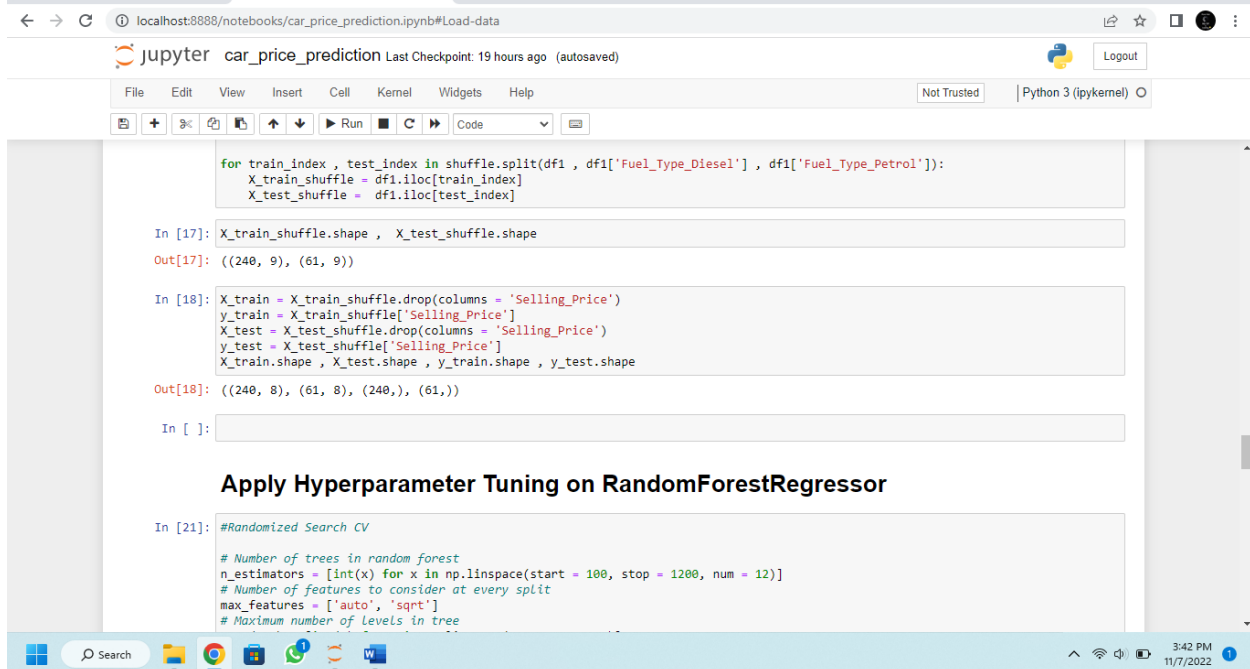
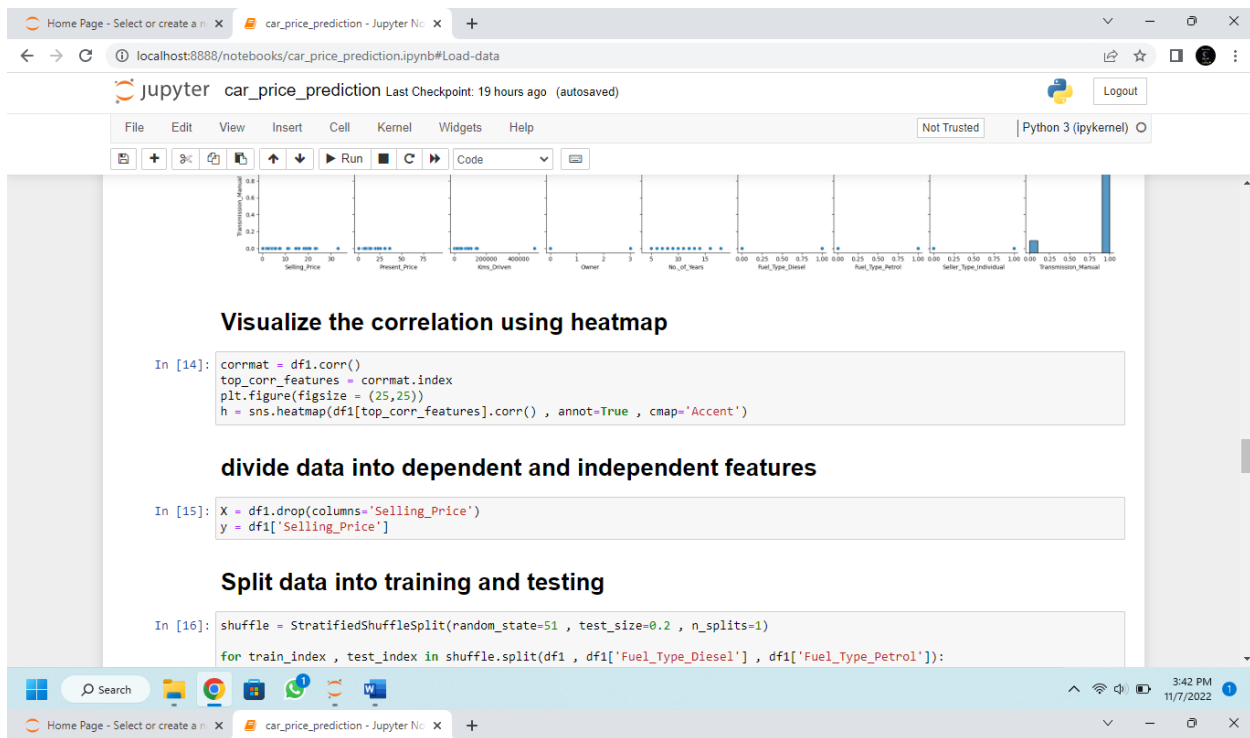


Model Building, Testing the Model, Model Evaluation





```
Home Page - Select or create a n... car_price_prediction - Jupyter No... +
localhost:8888/notebooks/car_price_prediction.ipynb#Load-data
jupyter car_price_prediction Last Checkpoint: 19 hours ago (autosaved) Logout
File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 (ipykernel) O
In [21]: #Randomized Search CV
# Number of trees in random forest
n_estimators = [int(x) for x in np.linspace(start = 100, stop = 1200, num = 12)]
# Number of features to consider at every split
max_features = ['auto', 'sqrt']
# Maximum number of levels in tree
max_depth = [int(x) for x in np.linspace(5, 30, num = 6)]
# max_depth.append(None)
# Minimum number of samples required to split a node
min_samples_split = [2, 5, 10, 15, 100]
# Minimum number of samples required at each leaf node
min_samples_leaf = [1, 2, 5, 10]

In [22]: # Create the random grid
random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf}

print(random_grid)

{'n_estimators': [100, 200, 300, 400, 500, 600, 700, 800, 900, 1000, 1100, 1200], 'max_features': ['auto', 'sqrt'], 'max_depth': [5, 10, 15, 20, 25, 30], 'min_samples_split': [2, 5, 10, 15, 100], 'min_samples_leaf': [1, 2, 5, 10]}

In [23]: rfr = RandomForestRegressor()
```

```
Home Page - Select or create a n... car_price_prediction - Jupyter No... +
localhost:8888/notebooks/car_price_prediction.ipynb#Load-data
jupyter car_price_prediction Last Checkpoint: 19 hours ago (autosaved) Logout
File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 (ipykernel) O
In [24]: # Random search of parameters, using 3 fold cross validation,
# search across 100 different combinations
rfr = RandomizedSearchCV(estimator = rfr, param_distributions = random_grid,scoring='neg_mean_squared_error',
                        n_iter = 10, cv = 5, verbose=2, random_state=51, n_jobs = 1)

In [ ]: rfr.fit(X_train,y_train)

Fitting 5 folds for each of 10 candidates, totalling 50 fits
[CV] END max_depth=25, max_features=auto, min_samples_leaf=10, min_samples_split=2, n_estimators=600; total time= 1.7s
[CV] END max_depth=25, max_features=auto, min_samples_leaf=10, min_samples_split=2, n_estimators=600; total time= 1.5s
[CV] END max_depth=25, max_features=auto, min_samples_leaf=10, min_samples_split=2, n_estimators=600; total time= 1.5s
[CV] END max_depth=25, max_features=auto, min_samples_leaf=10, min_samples_split=2, n_estimators=600; total time= 1.4s
[CV] END max_depth=25, max_features=auto, min_samples_leaf=10, min_samples_split=2, n_estimators=600; total time= 1.4s
[CV] END max_depth=10, max_features=auto, min_samples_leaf=1, min_samples_split=2, n_estimators=600; total time= 1.7s
[CV] END max_depth=10, max_features=auto, min_samples_leaf=1, min_samples_split=2, n_estimators=600; total time= 2.1s
[CV] END max_depth=10, max_features=auto, min_samples_leaf=1, min_samples_split=2, n_estimators=600; total time= 1.7s
[CV] END max_depth=10, max_features=auto, min_samples_leaf=1, min_samples_split=2, n_estimators=600; total time= 1.7s
[CV] END max_depth=10, max_features=auto, min_samples_leaf=1, min_samples_split=2, n_estimators=600; total time= 1.7s
[CV] END max_depth=20, max_features=sqrt, min_samples_leaf=2, min_samples_split=5, n_estimators=900; total time= 2.6s
[CV] END max_depth=20, max_features=sqrt, min_samples_leaf=2, min_samples_split=5, n_estimators=900; total time= 2.1s
[CV] END max_depth=20, max_features=sqrt, min_samples_leaf=2, min_samples_split=5, n_estimators=900; total time= 1.6s
[CV] END max_depth=20, max_features=sqrt, min_samples_leaf=2, min_samples_split=5, n_estimators=900; total time= 1.6s
[CV] END max_depth=20, max_features=sqrt, min_samples_leaf=1, min_samples_split=100, n_estimators=600; total time= 1.1s
[CV] END max_depth=20, max_features=sqrt, min_samples_leaf=1, min_samples_split=100, n_estimators=600; total time= 1.2s
[CV] END max_depth=20, max_features=sqrt, min_samples_leaf=1, min_samples_split=100, n_estimators=600; total time= 1.7s
```

Home Page - Select or create a notebook | car_price_prediction - Jupyter Notebook | +

localhost:8888/notebooks/car_price_prediction.ipynb#Load-data

jupyter car_price_prediction Last Checkpoint: 19 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 (ipykernel)

Predict data

```
In [ ]: pred = rfr.predict(X_test)
```

Check the r2_score

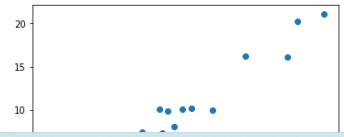
```
In [33]: r2_score(pred, y_test)
```

```
Out[33]: 0.9656759933799154
```

Visualize the actual and predicted data

```
In [25]: plt.scatter(y_test, pred)
```

```
Out[25]: <matplotlib.collections.PathCollection at 0x27310569c18>
```



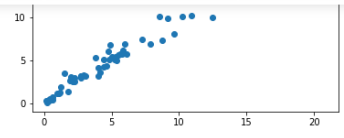
Windows taskbar: Search, File Explorer, Chrome, Jupyter Notebook, WhatsApp, Word. System tray: 3:43 PM, 11/7/2022.

Home Page - Select or create a notebook | car_price_prediction - Jupyter Notebook | +

localhost:8888/notebooks/car_price_prediction.ipynb#Load-data

jupyter car_price_prediction Last Checkpoint: 19 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 (ipykernel)



Display the actual and predicted data

```
In [26]: pd.DataFrame(np.c_[y_test, pred], columns=['Actual', 'Predicted'])
```

```
Out[26]:
```

	Actual	Predicted
0	12.50	10.006827
1	0.50	0.455760
2	0.65	0.705512
3	0.35	0.392024
4	14.90	16.263583
...
56	9.15	9.904631
57	4.75	6.072011
58	10.25	10.055072
59	0.38	0.313276

Windows taskbar: Search, File Explorer, Chrome, Jupyter Notebook, WhatsApp, Word. System tray: 3:43 PM, 11/7/2022.

Home Page - Select or create a notebook x car_price_prediction - Jupyter Notebook x +

localhost:8888/notebooks/car_price_prediction.ipynb#Load-data

jupyter car_price_prediction Last Checkpoint: 19 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 (ipykernel)

0 5 10 15 20

Display the actual and predicted data

```
In [26]: pd.DataFrame(np.c_[y_test , pred] , columns =['Actual' , 'Predicted'])
```

Out[26]:

	Actual	Predicted
0	12.50	10.006827
1	0.50	0.455760
2	0.65	0.705512
3	0.35	0.392024
4	14.90	16.263583
...
56	9.15	9.904631
57	4.75	6.072011
58	10.25	10.055072
59	0.38	0.313276
60	2.75	3.184455

61 rows x 2 columns

3:43 PM 11/7/2022